

# DevOps Interview Questions/Answers:

## Structure I'll follow:

- CI/CD Tools (Jenkins, GitHub Actions, GitLab CI/CD)
- Configuration Management (Ansible, Chef, Puppet)
- Containerization (Docker)
- Container Orchestration (Kubernetes, ECS, EKS)
- Infrastructure as Code (Terraform, CloudFormation)
- Monitoring and Logging (Prometheus, Grafana, ELK Stack)
- Cloud (AWS Services – EC2, S3, VPC, IAM, Lambda, RDS, ECS, EKS, etc.)
- Security and Networking in DevOps
- General DevOps Methodologies & Culture
- Scenario-Based Questions

## Entire guide now has 160+ questions with deep explanations.

### Section 1. CI/CD Tools (Jenkins, GitHub Actions, GitLab CI/CD)

#### Q1. What is CI/CD and why is it important in DevOps?

Answer:

- **CI (Continuous Integration)** is the practice of frequently integrating code into a shared repository and running automated tests to detect issues early.
- **CD (Continuous Delivery/Deployment)** ensures that the application is automatically deployed to staging or production environments after passing tests.
- **Importance:**
  - Reduces integration problems.
  - Speeds up release cycles.
  - Improves quality via automated testing.

Extra Tip: Mention tools like Jenkins, GitHub Actions, GitLab CI/CD, AWS CodePipeline and say which you've used in projects.

#### Q2. How does Jenkins work internally?

Answer:

- Jenkins is a **Java-based automation server** that runs on a servlet container like **Apache Tomcat**.
- It uses a **master-agent architecture**:

agents.

- **Agents (Nodes):** Execute build jobs.
- Jobs can be configured via **UI** or **Jenkinsfile (Pipeline as Code)**.
- Supports **plugins** for integration with tools like Git, Docker, Kubernetes, AWS.

**Extra Tip:** If asked about scaling Jenkins, talk about **Kubernetes-based Jenkins Agents** for elasticity.

### Q3. Difference between Declarative and Scripted Jenkins Pipelines?

**Answer:**

- **Declarative Pipeline:**

- Uses a structured, predefined syntax.
- Easier to read, more opinionated.

```
groovy

pipeline {
    agent any
    stages {
        stage('Build') { steps { echo 'Building...' } }
    }
}
```

### Scripted Pipeline:

- Written in Groovy, highly flexible.
- Good for complex logic.

```
groovy

node {
    stage('Build') {
        echo 'Building...'
    }
}
```

complex logic.

#### Q4. How do you create a Jenkins pipeline for a microservices app?

Answer:

- Create **Jenkinsfile** in repo:

```
groovy Copy Edit  
  
pipeline {  
    agent any  
    stages {  
        stage('Build') { steps { sh 'mvn clean package' } }  
        stage('Docker Build') { steps { sh 'docker build -t myapp .' } }  
        stage('Push to Registry') { steps { sh 'docker push myrepo/myapp' } }  
        stage('Deploy') { steps { sh 'kubectl apply -f k8s/deployment.yaml' } }  
    }  
}
```

- Configure **webhooks** for CI triggers.



#### Q5. What is the difference between Continuous Delivery and Continuous Deployment?

Answer:

- **Continuous Delivery:** Deployment to production requires **manual approval**.
- **Continuous Deployment:** Fully automated to production after tests pass.

#### Q6. How do you integrate Jenkins with GitHub?

Answer:

- Install **GitHub plugin** in Jenkins.
- Generate **Personal Access Token** in GitHub.
- Add webhook in GitHub pointing to <http://jenkins-server/github-webhook/>.
- Configure pipeline with SCM as GitHub repo.

#### Q7. How do you secure sensitive credentials in Jenkins?

Answer:

- Inject via withCredentials in pipeline:

```
groovy  
  
withCredentials([string(credentialsId: 'AWS_KEY', variable: 'AWS_KEY')]) {  
    sh "echo $AWS_KEY"  
}
```

#### Q8. Explain how you would set up a multi-branch pipeline in Jenkins.

Answer:

- Install **Multibranch Pipeline Plugin**.
- Configure Jenkins job to scan branches for Jenkinsfile.
- Each branch gets its own pipeline run automatically.

#### Q9. Difference between Jenkins and GitHub Actions?

Answer:

- **Jenkins**: Self-hosted, plugin-rich, more customizable.
- **GitHub Actions**: Cloud-based, built into GitHub, YAML workflows, auto-scaling runners.

#### Q10. How do you scale Jenkins for high availability?

Answer:

- Use **Jenkins master + multiple agents**.
- Run agents as **Kubernetes pods** using Jenkins Kubernetes Plugin.
- Use **Jenkins Operations Center** for distributed masters.

#### Q11. What is a Jenkins Shared Library?

Answer:

- A reusable pipeline code repository.
- Store common steps like build, test, deploy.
- Include in Jenkinsfile:

Follow [@coding.sight](#) And  
For Full Notes Comment “DevOps”



@coding.sight

```
@Library('my-shared-lib') _
```

## Q12. Explain Blue/Green Deployment in a Jenkins pipeline.

Answer:

- Deploy new version to **Green environment** while Blue is live.
- Test Green, then switch traffic.
- Rollback by pointing back to Blue if issues occur.

## Q13. What are GitHub Actions runners?

Answer:

- **Hosted runners:** Provided by GitHub (Linux, Windows, macOS).
- **Self-hosted runners:** Installed on your own infrastructure for customization.

## Q14. What is a GitLab CI/CD pipeline and how does it differ from Jenkins?

Answer:

- GitLab CI/CD uses .gitlab-ci.yml in the repo.
- Native to GitLab, integrated with repo & issues.
- Jenkins is external and integrates with multiple SCMs.

## Q15. How do you troubleshoot a failing Jenkins pipeline?

Answer:

- Check **console output logs**.
- Validate **syntax in Jenkinsfile** with lint.
- Enable **Pipeline Replay** for debugging.
- Check **agent connectivity** and plugin versions.

That's the **CI/CD section done with 15 detailed Q&A**.

## Section 2: Docker (Containerization)

### Q1. What is Docker and why do we use it?

Answer:

dependencies into a single **lightweight, portable container**.

- **Benefits:**

- Consistent environment across dev, test, prod.
- Faster deployments than VMs.
- Efficient resource utilization.

## Q2. Difference between Docker image and Docker container?

**Answer:**

- **Image:** Blueprint of the application with all dependencies. Immutable.
- **Container:** Running instance of an image.

## Q3. Explain Docker architecture.

**Answer:**

- **Client:** CLI or API that sends commands.
- **Docker Daemon:** Runs on the host, manages containers/images.
- **Registry:** Stores images (Docker Hub, AWS ECR).
- Communication via **REST API**.

## Q4. What is a Dockerfile?

**Answer:**

- A text file containing instructions to build a Docker image.
- Example:

```
dockerfile

FROM ubuntu:20.04
RUN apt-get update && apt-get install -y nginx
COPY index.html /var/www/html
CMD ["nginx", "-g", "daemon off;"]
```

## Q5. Explain Docker storage drivers.

**Answer:**

- **Overlay2 (Linux):** Default, efficient copy-on-write.

- **aufs, devicemapper:** Older alternatives.
- Allows container layers to share base image layers.

## Q6. How does Docker networking work?

Answer:

- **Bridge (default):** Containers communicate via internal network.
- **Host:** Shares host network stack.
- **Overlay:** For multi-host networking (Swarm/Kubernetes).
- **None:** No networking.

## Q7. What is the difference between ENTRYPPOINT and CMD in Dockerfile?

Answer:

- **ENTRYPOINT:** Always executed as the main process.
- **CMD:** Provides default arguments to ENTRYPOINT.
- Example:

```
dockerfile
```

```
ENTRYPOINT ["python"]
CMD ["app.py"]
```

## Q8. How do you reduce Docker image size?

Answer:

- Use smaller base images (e.g., **alpine**).
- Combine RUN commands with && to reduce layers.
- Remove unnecessary packages and temp files.
- Use **multi-stage builds**.

## Q9. Explain Docker Compose.

Answer:

- Tool to define and manage multi-container apps using docker-compose.yml.
- Example:

Follow [coding.sight](https://@coding.sight) And  
For Full Notes Comment “**DevOps**”   

@coding.sight

```
version: '3'  
services:  
  web:  
    image: nginx  
  db:  
    image: mysql
```

## Q10. How do you secure Docker containers?

Answer:

- Use **non-root users**.
- Keep images updated.
- Use **Docker Bench for Security**.
- Scan images with **Trivy or Clair**.

## Q11. What is Docker Swarm and how does it compare to Kubernetes?

Answer:

- Docker Swarm: Native clustering tool for Docker.
- Easier setup than Kubernetes, but less feature-rich.
- Kubernetes is the industry standard now.

## Q12. How do you persist data in Docker?

Answer:

- Use **Volumes** (docker volume create) or **Bind mounts**.
- Example:

```
bash  
  
docker run -v /host/path:/container/path myapp
```

## Q13. How do you troubleshoot a Docker container that keeps crashing?

Answer:

- Inspect container: docker inspect.
- Check resource limits and entrypoint errors.

#### Q14. What is the difference between Docker and a Virtual Machine?

Answer:

- **Docker:** Uses host OS kernel, lightweight.
- **VM:** Requires full OS per instance, heavy.
- Containers start in seconds; VMs take minutes.

#### Q15. What is a multi-stage build in Docker?

Answer:

- Allows using multiple FROM statements to reduce final image size.
- Example:

```
dockerfile

FROM golang:1.18 AS builder
WORKDIR /app
RUN go build -o app

FROM alpine:latest
COPY --from=builder /app/app .
CMD ["./app"]
```

 Docker section complete.

### SECTION 3: Kubernetes (Container Orchestration)

#### Q1. What is Kubernetes and why do we use it?

Answer:

Kubernetes (K8s) is an **open-source container orchestration platform** for automating deployment, scaling, and management of containerized applications.

- Features: **Self-healing, Auto-scaling, Service discovery, Rolling updates.**

#### Q2. Explain Kubernetes architecture.

[@coding.sight](https://www.coding.sight)

- **Master Node (Control Plane):**
  - API Server (handles requests), Scheduler (assigns pods), Controller Manager, etcd (key-value store).
- **Worker Nodes:**
  - Kubelet (manages pods), Kube-proxy (networking), Container Runtime (Docker/Containerd).

### Q3. What is a Pod in Kubernetes?

**Answer:**

- Smallest deployable unit in K8s.
- Can contain **one or more containers** sharing the same network namespace and storage volumes.

### Q4. Difference between ReplicaSet and Deployment?

**Answer:**

- **ReplicaSet:** Ensures a certain number of pod replicas run at any time.
- **Deployment:** Manages ReplicaSets and enables **rolling updates & rollbacks**.

### Q5. Explain Kubernetes Services.

**Answer:**

- **ClusterIP (default):** Internal communication.
- **NodePort:** Exposes service on node's IP at a static port.
- **LoadBalancer:** Uses cloud provider's LB for external access.
- **ExternalName:** Maps to an external DNS name.

### Q6. What are ConfigMaps and Secrets?

**Answer:**

- **ConfigMap:** Stores non-sensitive config data (e.g., environment variables).
- **Secret:** Stores sensitive data (e.g., passwords, API keys). Encoded in Base64.

### Q7. Explain Kubernetes Ingress.

**Answer:**

- Ingress is an API object that **manages external access** (HTTP/HTTPS) to services in a cluster.

## **Q8. How does Kubernetes handle service discovery?**

**Answer:**

- Each service gets a **DNS name** via CoreDNS.
- Pods can communicate using service name instead of IP.

## **Q9. Explain Rolling Update vs Recreate Strategy.**

**Answer:**

- **RollingUpdate (default):** Updates pods incrementally without downtime.
- **Recreate:** Deletes all old pods, then creates new ones (downtime possible).

## **Q10. How does Kubernetes handle scaling?**

**Answer:**

- **Horizontal Pod Autoscaler (HPA):** Scales pods based on CPU/memory metrics.
- **Vertical Pod Autoscaler (VPA):** Adjusts pod resources.
- **Cluster Autoscaler:** Adds/removes worker nodes.

## **Q11. What is etcd in Kubernetes?**

**Answer:**

- A distributed **key-value store** that stores all cluster state data.
- Highly available and fault-tolerant.

## **Q12. Explain taints and tolerations.**

**Answer:**

- **Taint:** Applied to nodes to restrict which pods can run on them.
- **Toleration:** Allows pods to run on tainted nodes.

## **Q13. How do you monitor Kubernetes clusters?**

**Answer:**

- Use **Prometheus + Grafana, Kube-State-Metrics, ELK stack, Lens dashboard.**

## **Q14. How does Kubernetes manage networking?**

- Flat network model: Every pod gets a unique IP.
- Uses **CNI plugins** like Flannel, Calico, Weave for networking.

#### **Q15. What is a StatefulSet in Kubernetes?**

**Answer:**

- Manages stateful applications.
- Provides **stable network identities** and **persistent storage** for each pod.

#### **Q16. Explain Kubernetes Namespaces.**

**Answer:**

- Logical partitions for organizing resources.
- Useful for multi-team clusters and resource isolation.

#### **Q17. What is the difference between Docker Swarm and Kubernetes?**

**Answer:**

- Swarm: Simple, less powerful, native to Docker.
- Kubernetes: Feature-rich, industry standard for orchestration.

#### **Q18. How do you troubleshoot a failing pod?**

**Answer:**

- Check logs: kubectl logs pod-name.
- Describe pod: kubectl describe pod pod-name.
- Events: kubectl get events.

#### **Q19. What is the difference between NodePort and LoadBalancer?**

**Answer:**

- **NodePort:** Exposes service on all nodes at a fixed port (manual LB).
- **LoadBalancer:** Uses cloud LB for automatic external access.

#### **Q20. How do you secure Kubernetes clusters?**

**Answer:**

- Enable **RBAC**.
- Use **Network Policies**.

- Limit container privileges.

 **Kubernetes section complete.**

## SECTION 4: Infrastructure as Code (Terraform & CloudFormation)

### Q1. What is Infrastructure as Code (IaC)?

**Answer:**

- IaC is the practice of managing infrastructure using **code**, not manual processes.
- Benefits: **Consistency, Version Control, Automation.**

### Q2. What is Terraform and why is it popular?

**Answer:**

- Open-source IaC tool by HashiCorp.
- Uses **HCL (HashiCorp Configuration Language)**.
- Multi-cloud support (AWS, Azure, GCP).

### Q3. Explain Terraform architecture.

**Answer:**

- **Core:** Reads config and manages lifecycle.
- **Providers:** Plugins for cloud platforms.
- **State File:** Stores infrastructure state for tracking changes.

### Q4. What is the difference between Terraform and CloudFormation?

**Answer:**

- **Terraform:** Multi-cloud, uses HCL.
- **CloudFormation:** AWS-only, uses JSON/YAML.

### Q5. What is Terraform state and why is it important?

**Answer:**

- State file (`terraform.tfstate`) tracks resources created by Terraform.
- Enables incremental updates.
- Should be stored in **remote backends** (S3, etc.).

## Q6. What are Terraform backends?

**Answer:**

- Mechanism for storing state remotely (S3, GCS, Consul).
- Allows **collaboration** and **locking**.

## Q7. Explain terraform init, plan, apply, destroy.

**Answer:**

- init – Initialize working directory.
- plan – Show execution plan.
- apply – Apply changes.
- destroy – Delete resources.

## Q8. How do you manage secrets in Terraform?

**Answer:**

- Use **Vault**, **AWS Secrets Manager**, or environment variables.
- Never hardcode secrets in .tf files.

## Q9. Explain Terraform modules.

**Answer:**

- Reusable configurations stored in separate directories or repos.
- Called using module block in main .tf file.

## Q10. How do you handle drift in Terraform?

**Answer:**

- Use terraform plan to detect differences between real infrastructure and state file.
- Run terraform apply to fix.

## Q11. What is a Terraform provider?

**Answer:**

- Plugin that interacts with APIs (AWS, Azure, GCP).
- Declared in configuration:

Follow **@coding.sight** And  
For Full Notes Comment "DevOps"



**@coding.sight**

hcl

```
provider "aws" {  
    region = "us-east-1"  
}
```

## Q12. How do you upgrade Terraform version safely?

**Answer:**

- Use `terraform version` to check.
- Run `terraform O.x upgrade` command if major changes.

## Q13. What is the difference between `count` and `for_each` in Terraform?

**Answer:**

- **count:** Simple replication based on an integer.
- **for\_each:** Creates resources based on a map or set.

## Q14. Explain `depends_on` in Terraform.

**Answer:**

- Forces explicit resource dependency.
- Useful when implicit dependencies are not enough.

## Q15. How do you ensure team collaboration in Terraform?

**Answer:**

- Use **remote backend with locking** (e.g., S3 + DynamoDB).
- Use **Terraform Cloud/Enterprise** for governance.

 **Terraform section done.**

## SECTION 5: Configuration Management (Ansible, Chef, Puppet)

### Q1. What is configuration management in DevOps?

**Answer:**

Configuration management is the process of **automating the setup and**

environments.

## Q2. What is Ansible and why is it popular?

Answer:

- **Agentless**, uses **SSH** for communication.
- Written in Python, uses **YAML playbooks**.
- Easy to learn, no extra software on nodes.

## Q3. Explain Ansible architecture.

Answer:

- **Control Node**: Runs playbooks.
- **Managed Nodes**: Machines being configured.
- **Inventory**: List of hosts.
- **Modules**: Pre-built scripts for tasks.

## Q4. What is an Ansible playbook?

Answer:

- A YAML file that defines **tasks to execute on target hosts**.

```
yaml

- hosts: webservers
  tasks:
    - name: Install NGINX
      apt:
        name: nginx
        state: present
```

## Q5. How does Ansible differ from Chef and Puppet?

Answer:

- **Ansible**: Agentless, YAML, simple.
- **Chef**: Uses Ruby DSL, requires agent.
- **Puppet**: Declarative, uses manifests, agent-based.

**Answer:**

- Use **Ansible Vault** to encrypt sensitive data:

**Q7. What are Ansible roles?**

**Answer:**

- A way to organize playbooks into reusable components (tasks, vars, handlers).

**Q8. How do you test Ansible playbooks before running in production?**

**Answer:**

- Use --check flag for dry-run.
- Use **Molecule** for testing.

**Q9. How does Ansible ensure idempotency?**

**Answer:**

- Tasks are written to produce the same result even if executed multiple times.

**Q10. What is the difference between Ansible ad-hoc commands and playbooks?**

**Answer:**

- **Ad-hoc:** One-time tasks.
- **Playbooks:** Structured, reusable automation.

**Q11. How do you handle dynamic inventories in Ansible?**

**Answer:**

- Use scripts or **cloud inventory plugins** (AWS, GCP).

**Q12. What are Ansible facts?**

**Answer:**

- System properties gathered by Ansible before running tasks.

**Q13. How do you scale Ansible for large environments?**

- Use **Ansible Tower/AWX** for GUI and RBAC.

#### **Q14. What is the difference between handlers and tasks in Ansible?**

**Answer:**

- **Handlers:** Run only when notified (e.g., restart service).
- **Tasks:** Run in order always.

#### **Q15. How do you debug Ansible playbooks?**

**Answer:**

- Use -vvvv for verbose mode.
- Use debug module to print variables.

 **Ansible section complete.**

### **SECTION 6: Monitoring & Logging (Prometheus, Grafana, ELK)**

#### **Q1. Why is monitoring important in DevOps?**

**Answer:**

- Ensures system reliability, detects issues early, supports **SLO/SLI/SLA compliance**.

#### **Q2. What is Prometheus and why is it used?**

**Answer:**

- Open-source **metrics-based monitoring** system.
- Pulls metrics from targets via HTTP.
- Stores data in **time-series database**.

#### **Q3. How does Prometheus architecture work?**

**Answer:**

- **Server:** Scrapes and stores metrics.
- **Exporters:** Collect metrics (e.g., node\_exporter).
- **Alertmanager:** Sends alerts.

#### **Q4. What is Grafana?**

- Visualization tool for Prometheus, InfluxDB, Elasticsearch.
- Creates dashboards and alerts.

## Q5. Explain ELK Stack.

Answer:

- **Elasticsearch:** Stores logs.
- **Logstash:** Collects & processes logs.
- **Kibana:** Visualization UI.

## Q6. Difference between push and pull metrics in Prometheus?

Answer:

- **Pull:** Prometheus scrapes targets.
- **Push:** Targets send metrics to Pushgateway (rare case).

## Q7. How do you monitor Kubernetes clusters?

Answer:

- Prometheus + Grafana with **kube-state-metrics** and **cAdvisor**.

## Q8. What is a Prometheus Alert Rule?

Answer:

- Defines conditions for alerts.

```
yaml
  - alert: HighCPU
    expr: node_cpu_seconds_total > 90
    for: 5m
```

## Q9. How do you secure Prometheus and Grafana?

Answer:

- Use RBAC in Grafana.

#### Q10. What are blackbox and whitebox monitoring?

Answer:

- **Blackbox:** Tests from outside (ping, HTTP).
- **Whitebox:** Internal metrics.

#### Q11. How do you monitor logs in microservices?

Answer:

- Use **centralized logging** with ELK or Loki.

#### Q12. How do you set up alerts in Grafana?

Answer:

- Create thresholds on panels, integrate with Slack, email, or PagerDuty.

 **Monitoring section complete.**

### SECTION 7: AWS Cloud Services

#### Q1. What are the core AWS services for DevOps?

Answer:

- **Compute:** EC2, Lambda, ECS, EKS.
- **Storage:** S3, EBS, EFS.
- **Networking:** VPC, Route 53, ALB/NLB.
- **CI/CD:** CodeBuild, CodeDeploy, CodePipeline.
- **Monitoring:** CloudWatch, X-Ray.

#### Q2. Explain EC2 instance types.

Answer:

- **General Purpose (t2, t3):** Balanced.
- **Compute Optimized (c5):** High CPU.
- **Memory Optimized (r5):** High RAM.
- **GPU (p2, g4):** ML, graphics.

#### Q3. What is an Auto Scaling Group (ASG)?

Follow @coding.sight And  
For Full Notes Comment "DevOps"    ... 